

Assignment 1

Shayan Shabanzadeh

Experiment Design

We measured how parallelism affects both (i) **runtime** and (ii) **accuracy** when approximating π using three implementations:

- **Integration (atomic):** trapezoidal rule with a shared accumulator updated via *atomic* operations.
- **Integration (critical):** same, but the accumulator update is protected by a *critical* section.
- **Monte Carlo (MC):** throw random points and count hits inside the unit quarter-circle.

For each method we varied:

- **Threads:** $T \in \{1, 2, 4, 8, 16, 32, 64, 128\}$.
- **Work budget:** steps (integration) or guesses (MC) in $\{10^5, 10^6, 10^7, 10^8\}$.

For every configuration we computed:

$$\text{Runtime} = \text{mean wall time (s)}, \quad \text{MAE} = |\hat{\pi} - \pi|.$$

Figures use actual thread counts on an x-axis with \log_2 scaling for readability; the data points are the real thread counts.

Results: Runtime (What We Saw)

MC scales; integration does not. The *single* most important observation is that MC speedups are substantial with more threads, while both integration variants slow down due to per-iteration sharing.

Monte Carlo (MC). For the largest budget (guesses = 10^8), mean time falls from **1.3701 s** (1 thread) to **0.03823 s** (128 threads) — about **35.8× speedup**. Speedups are visible at all budgets:

- 10^5 : 0.001 349 s \rightarrow 0.001 143 s (1 \rightarrow 128); small absolute gain because the job is tiny.

- 10^6 : 0.012 968 s \rightarrow 0.000 328 s (1 \rightarrow 128); \sim **39.6×**.
- 10^7 : 0.134 291 s \rightarrow 0.017 369 s (1 \rightarrow 128); \sim **7.7×**.
- 10^8 : 1.370 114 s \rightarrow 0.038 228 s (1 \rightarrow 128); \sim **35.8×**.

Why: MC trials are independent (no shared updates), so threads progress without blocking.

Integration (atomic). Here, every step updates a shared accumulator via an atomic. Contention grows with T ; stalls dominate for large budgets. For steps = 10^8 , runtime *increases* from **0.6857 s** (1 thread) to **13.5628 s** (128 threads): a **19.8× slowdown**. Similar trend at other budgets:

- 10^5 : 0.000 784 s (1) \rightarrow 0.021 986 s (128) (*overhead dominates when work is tiny*).
- 10^6 : 0.006 928 s (1) \rightarrow 0.132 540 s (128) \uparrow .
- 10^7 : 0.068 570 s (1) \rightarrow 1.344 519 s (128) \uparrow .
- 10^8 : 0.685 736 s (1) \rightarrow 13.562 825 s (128) \uparrow .

Integration (critical). The critical section forces *full* serialization at the update, so it behaves even worse than atomics. For steps = 10^8 , time grows from **1.0561 s** (1) to **39.8585 s** (128): a **37.8× slowdown**. Other budgets show the same pattern:

- 10^5 : 0.001 126 s (1) \rightarrow 0.047 521 s (128).
- 10^6 : 0.010 311 s (1) \rightarrow 0.434 496 s (128).
- 10^7 : 0.104 327 s (1) \rightarrow 4.099 531 s (128).
- 10^8 : 1.056 109 s (1) \rightarrow 39.858 507 s (128).

Summary of runtime. MC: *parallel speedups* (best at large budgets). Integration with per-iteration sharing: *parallel slowdowns* that worsen with T and budget. The runtime plot (left) visually summarizes these numeric trends.

Results: Accuracy (What We Saw)

Accuracy is governed by the work budget, not the thread count. Increasing steps/guesses from 10^5 to 10^8 consistently reduces MAE by 1–2 orders of magnitude. For a *fixed* budget, errors fluctuate slightly with T but show *no systematic dependence* on threads.

Monte Carlo (MC). At 1 thread, MAE improves from **1.678e-3** (10^5) to **2.033e-5** (10^8), matching the familiar $O(1/\sqrt{N})$ behavior. Holding budget fixed at 10^8 , MAE across threads stays in the same $O(10^{-5})$ band:

- 10^8 guesses: $2.033 \cdot 10^{-5}$ (1) vs. $3.927 \cdot 10^{-5}$ (128).

Integration (atomic). At 1 thread, MAE shrinks from **1.415e-3** (10^5) to **1.831e-5** (10^8). For a fixed budget of 10^8 , MAE remains $O(10^{-5})$ across threads:

- 10^8 steps: $1.831 \cdot 10^{-5}$ (1) vs. $4.062 \cdot 10^{-5}$ (128).

Integration (critical). At 1 thread, MAE goes from **1.399e-3** (10^5) to **1.674e-5** (10^8). For 10^8 steps across threads, it stays $O(10^{-5})$:

- 10^8 steps: $1.674 \cdot 10^{-5}$ (1) vs. $5.308 \cdot 10^{-5}$ (128).

Summary of accuracy. More samples \Rightarrow better accuracy. Threads simply change how *fast* we can reach a given budget (and for integration with sharing, they actually make it slower); they do not directly change the error for a fixed budget. The error plot (right) shows essentially flat curves vs. thread count for each budget.

Why These Results Happen (Mechanistic Explanation)

MC is embarrassingly parallel: each trial is independent; no shared state \Rightarrow no contention. *Integration with per-iteration sharing:* every step touches a single shared accumulator. With **atomics**, the cache line and RMW pipeline serialize many updates; with a **critical** section, only one thread can update at a time, producing the strongest slowdown. Hence:

MC: scale with T vs. Integration: stall with T .

Accuracy follows the mathematics of the estimator (MC variance $\propto 1/N$; trapezoidal error decreases with finer discretization). Threads do not change that estimator, only the throughput of samples.




Figure 1: Runtime vs. threads (\log_2 ticks). MC accelerates with threads (e.g., 10^8 guesses: $1.37\text{ s} \rightarrow 0.038\text{ s}$); integration with atomics/critical slows down strongly as contention rises (e.g., 10^8 steps atomic: $0.686\text{ s} \rightarrow 13.56\text{ s}$; critical: $1.06\text{ s} \rightarrow 39.86\text{ s}$).




Figure 2: Mean absolute error (MAE) vs. threads. Curves are essentially flat for a fixed budget, confirming error depends on total samples, not thread count (e.g., MC 10^8 : 2.0×10^{-5} – 3.9×10^{-5} across T).