

# GPU SPMV Performance Evaluation: CSR vs. ELL

## 1. Overview

This report evaluates two GPU-accelerated Sparse Matrix–Vector Multiplication (SpMV) implementations:

- CSR-based SpMV (row-compressed format)
- ELL-based SpMV (Ellpack format)

Both implementations were integrated into the provided CUDA codebase, and the missing functions were fully implemented:

1. convert\_csr\_to\_ell()
2. allocate\_csr\_gpu()
3. allocate\_ell\_gpu()
4. spmv\_kernel() (CSR kernel)
5. spmv\_kernel\_ell() (ELL kernel)

The matrix used was `cant.mtx`, and all experiments were executed with `MAX_ITER = 100` iterations for stable timing.

**CUDA version note:** The Talapas cluster required switching CUDA from the default version to **CUDA 11.2.0**. Because of this, we updated the environment setup in `spmv.srun`.

This resolved compatibility issues and ensured that both CSR and ELL kernels compiled correctly.

## 2. Implementation Summary

### 2.1 CSR GPU Implementation

The CSR GPU kernel assigns one thread block per matrix row:

- Threads iterate over the nonzeros of the row using a strided loop.
- Partial sums are stored in `sdata[]` in shared memory.
- A parallel reduction computes the final row result.

#### GPU execution parameters:

```
blocks = GRID_SIZE
threads = BLOCK_SIZE
shared = threads * sizeof(double)
```

### 2.2 ELL GPU Implementation

The ELL format stores exactly `max_nnz_per_row` values per row. Each GPU block computes a single row:

- Threads iterate across the padded column slots.
- Zeros represent padded entries and are skipped.
- Shared-memory reduction is performed identically to CSR.

Because ELL uses column-major storage (`index = k*m + row`), memory accesses are more regular but sometimes include padded zeros.

### 2.3 Converting CSR to ELL

We implement:

1. Compute the maximum nonzeros in any row
2. Allocate  $m * \text{max\_nnz}$  slots
3. Copy each row's columns/values into ELL's padded layout

This ensures GPU-friendly, uniform row lengths.

### 2.4 GPU Memory Allocation

Both `allocate_csr_gpu` and `allocate_ell_gpu` use:

- `cudaMallocHost()` for pinned CPU memory
- `cudaMemcpy()` for fast transfer to device

This significantly improves PCIe bandwidth.

## 3. Experiment Design

We varied two execution parameters passed via `GRID_SIZE` and `BLOCK_SIZE`:

$$\text{Block sizes} = \{32, 64, 128, 256, 512\}$$

$$\text{Grid sizes} = \{3903, 7806, 15612, 31225, 62451\}$$

These correspond to:

- different numbers of thread blocks,
- different numbers of GPU threads per block.

For each configuration, we recorded:

- CSR execution time
- ELL execution time
- Speedup: ELL/CSR

## 4. Results and Analysis

### 4.1 Execution Time Differences

ELL is consistently slower than CSR across all parameter combinations. This is expected because:

1. ELL includes padded entries, increasing total operations
2. CSR has perfectly compact nonzero storage
3. CSR memory accesses are less wasteful




Figure 1: Execution time difference (ELL – CSR).

### 4.2 CSR vs ELL per Grid Size

Execution times rise smoothly with block size. CSR remains faster in every case, especially for large matrices.




Figure 2: CSR vs ELL execution time for each grid size.

### 4.3 Speedup Behavior

Speedup is defined as:

$$\text{Speedup} = \frac{\text{ELL time}}{\text{CSR time}}$$

Values above 1.0 mean CSR is faster.




Figure 3: Speedup (ELL / CSR). Lower is better for ELL.

Observations:

- ELL is worst at small block sizes (32–64)
- Speedup peaks near 64 threads per block
- For large block sizes (256–512), ELL performance declines sharply

This behavior matches ELL's sensitivity to padded elements.

## 5. Conclusions

- CSR clearly outperforms ELL for SpMV on GPUs in all tests.
- ELL suffers from padded zero entries and performs more work.
- CSR scales smoothly with block size and grid size.
- ELL does not gain performance from larger blocks; padding overhead grows.
- Switching CUDA to 11.2.0 was necessary for successful compilation and execution.

Overall, CSR is the preferred GPU format for sparse matrix–vector multiplication in this assignment.